

Michael Gardner

Information Extraction

Assignment 3 – Relation Classification

Runtime parameters for all reported scores:

Train file: semevalTrainReal.tsv

Dev file: semevalDev.tsv

Lr: 0.00005

Batch size: 16

Max epochs: 20

Implemented cmd line additions:

Llm : true or false

Llm\_choice: "distilbert-base-uncased", "bert-base-uncased", "openai-gpt"

Test best: whether to load model ('best\_performers/CNN\_bert-base-uncased')

This assignment, I built upon the models from previous assignments and presumed I would use the LSTM, which was the best performer there. The simple DANClassifier model which came in the starter code, with no alterations achieved a benchmark around 25% val accuracy and it was a bare minimum to beat. The base bidirectional LSTM came in with 23%. The experiment progressed from there to involve the Truncation.

Truncation is to say, the target sentence is shaved down to (entity1:entity2), inclusive so we can gain some knowledge from the target entity words themselves. The base LSTM improved with truncation, and we reached 37% accuracy and 0.27 F1. Other experiments with the entity-tags functionality didn't show any improvement over truncation, so it was shelved.

So the next upgrade was to stop using random vocab indexes and move on to embeddings. The code allows for llm to be turned on and llm\_choice to be selected from distilbert, base bert, and gpt. Checkpointed models could also be stored to save their best performing valid\_acc epoch.

**ALL SCORES REFLECT THE DEV or “VALIDATION” SET.**

	Val_acc	macro_f1	train_loss	sec/epoch
LSTM Truncated:	0.373	0.270	1.0	2.5s
Distilbert LSTM Truncated:	0.621	0.600	0.68	18s
BERTbase LSTM Truncated:	0.621	0.601	0.56	26s
GPT LSTM Truncated:	0.661	0.667	0.25	25s

Validation officially passed 60% baseline! And GPT is running away with it!

But how do we improve from here. My hypothesis in the model proposal was that truncation would show improvement and we've born that out at the base LSTM... but what about with embeddings. I needed to check.

	Val_acc	macro_f1	train_loss	sec/epoch
LSTM:	0.229	0.067	2.37	3s
Distilbert LSTM:	0.715	0.709	0.07	20s
BERTbase LSTM:	0.740	0.745	0.10	31s
GPT LSTM:	0.695	0.696	0.13	33s

That blew all my expectations away. GPT with full context loses accuracy, but BERT gains? Surpassed the 72% mark? If that's true, maybe we need the CNN from assignment 2; would that improve things?

The implementation of the CNN and LSTM from A2 each did take some alterations to accept this data, including excluding the masking functionality in favor the variable padding implemented in dataset, and the dataset code had to handle "input\_ids" separately from tokens.

So, the CNN, using 100 filters of widths 4,3,2...

	Val_acc	macro_f1	train_loss	sec/epoch
Distilbert CNN:	0.7275	0.7323	0.0468	20s
BERTbase CNN:	0.7450	0.7444	0.0270	34s
GPT CNN:	0.7090	0.7064	0.885	33s

BERT embeddings shine once more. That best performing model (epoch 13) can be loaded with the command line command `-test-best`. I'd call those BERT results for the LSTM and CNN within a statistical margin of error, which was also the case in A2, neck-and-neck. But the distilbert version saw moderate improvement, so we can say the CNN makes gains over the LSTM, no matter the embedding method.